

When Positive-Feedback Strategies Crowd: A Family-Level Threshold Framework via Agent-Based Simulation*

Yi-Hao Lai[†] VolPred Research System

April 2026

Abstract

Volatility targeting (VT), trend-following (TF), and short-horizon mean-reversion (MR) share a common procyclical structure: each strategy’s position update is a non-trivial function of recent prices or volatility, so widespread adoption can in principle generate self-reinforcing trading flows. Whether the resulting crowding threshold is VT-specific or a generic property of positive-feedback strategies is unsettled. We build an agent-based model with 1,000 heterogeneous agents, a Kyle (1985) market maker, endogenous VIX dynamics, and 200 fixed noise traders, and use it to test the threshold structure across three strategy classes under matched microstructure. Across 43,300 Monte Carlo simulations spanning 7 adoption levels $\times$ 4 treatments $\times$ 12 strategy-spec cells $\times$ 5 (λ, γ) OAT cells, three findings emerge. First, under a Sharpe-only detector calibrated to the standalone VT benchmark, VT crosses at 70% adoption—exactly the threshold reported by the canonical analysis. Second, TF and MR cross *at or below* the VT threshold in 12/12 strategy-spec cells and 5/5 microstructure cells, establishing positive-feedback crowding as

*We thank the VolPred Research System for computational support and OpenAI Codex / Claude code-reviewer agents for adversarial code review. Simulation code and replication data available upon request.

[†]Department of Finance, Da-Yeh University. Corresponding author. Email: yhlai@mail.dyu.edu.tw

a *family-level* phenomenon. Third, the qualitative ordering survives $\pm 50\%$ perturbations of Kyle- λ and VIX-feedback γ ; only the VT magnitude shifts directionally with λ , consistent with the Kyle-impact mechanism rather than a knife-edge artifact. VT remains the empirically dominant case (largest assets-under-management, real-world deployment), but the threshold is generic.

Keywords: volatility targeting, positive feedback, crowding, agent-based model, tipping point, systemic risk

JEL Classification: G11, G12, G17, G18

1 Introduction

Volatility targeting (VT) has become a ubiquitous tool in dynamic portfolio management. The canonical framework, formalized by [Moreira and Muir \[2017\]](#), prescribes scaling equity exposure as $w_t = \sigma^*/\hat{\sigma}_t$, where σ^* is the target volatility and $\hat{\sigma}_t$ is a conditional volatility estimate. [Harvey et al. \[2018\]](#) demonstrate that such strategies meaningfully reduce tail risk across asset classes, and the approach has been widely adopted by pension funds, insurance companies, and target-date funds. Industry estimates suggest that short-volatility and volatility-sensitive strategies collectively manage over USD 2 trillion in assets [[Cole, 2017](#)]. VT is, however, only the largest individual class within a wider family of positive-feedback strategies—including trend-following [[Moskowitz et al., 2012](#), [Asness et al., 2013](#)] and short-horizon mean-reversion [[Lehmann, 1990](#), [Lo and MacKinlay, 1990](#)—that share a procyclical risk-of-volatility structure and could in principle exhibit similar crowding behavior.

Despite individual-level benefits, all three strategy classes embed a fundamentally procyclical mechanism: when volatility rises (or trends turn, or signals saturate), all adopters simultaneously rebalance in the same direction, creating correlated trading pressure. The crowding risk literature warns of such procyclical dynamics—[Baltas \[2019\]](#) documents crowding effects in alternative risk premia, the European Central Bank [ECB, [European Central Bank, 2020](#)] explicitly flagged VT procyclicality as a source of market fragility, and [Cont and Bouchaud \[2000\]](#) build an agent-based herd-behavior model in which correlated demand generates fat-tailed returns—but no study provides a *strategy-class-by-strategy-class* quantitative threshold. Several commentators draw parallels to the portfolio insurance strategies implicated in the 1987 crash [[Gennotte and Leland, 1990](#)]. [Bookstaber et al. \[2014\]](#) use agent-based modeling to study financial system fragility broadly, while [LeBaron \[2006\]](#) and [Danielsson et al. \[2012\]](#) study feedback-driven market dynamics; applications that *explicitly compare* VT crowding against TF or MR crowding under identical microstructure are absent.

A critical gap remains: *at what adoption level does correlated positive-feedback trading become destabilizing, and is the threshold strategy-specific or a generic property of the*

feedback class? Existing analyses are either qualitative [European Central Bank, 2020], focused on a single strategy [Baltas, 2019], or lack a quantitative threshold altogether. The empirical VT-alpha literature itself is contested—Cederburg et al. [2020] and Liu et al. [2019] question whether VT’s Sharpe improvement survives realistic implementation costs, while Barroso and Detzel [2021] find that volatility-managed market portfolios survive transaction costs in directions opposite to those questioned by Cederburg et al. [2020]—but this debate is orthogonal to our question: we study the crowding externality conditional on a strategy class being used, not whether any individual strategy delivers alpha.

We address this gap with an agent-based simulation in which the same Kyle-microstructure (constant λ , endogenous VIX) is exposed to four interchangeable agent treatments: VT, TF, MR, and a NoiseControl falsifiability anchor. Our contribution is threefold. First, we deliver a *strategy-family threshold framework* in which the canonical VT 70% threshold reported by the standalone analysis is reproduced *exactly* under the same microstructure when a Sharpe-only detector is applied to the cell1 baseline ($\lambda = 0.005$, $\gamma = 200$, see §3.6). Second, we document a *cross-strategy threshold ordering*: TF and MR thresholds are at or below VT in 12/12 strategy-spec cells (Phase 2: scaling $\times$ window robustness, §3.4) and in 5/5 OAT cells ($\lambda/\gamma \pm 50\%$ robustness, §3.6), establishing positive-feedback crowding as a family-level rather than VT-specific phenomenon. Third, we provide *joint robustness*: the qualitative ordering $\text{TF/MR} \leq \text{VT}$ survives both strategy-spec perturbation and microstructure perturbation; only the VT magnitude shifts directionally with the Kyle parameter λ , consistent with the impact mechanism encoded in Eq. (4) rather than a knife-edge artifact. The fixed-vs-scaled-liquidity validation that featured prominently in earlier work is retained as a methodological check (§3.7) and confirms that approximately half of the originally observed degradation is attributable to liquidity evaporation rather than crowding per se. These results provide a first quantitative cross-strategy benchmark for assessing positive-feedback crowding risk.

2 Model

We construct an agent-based model with $N = 1,000$ heterogeneous agents operating in a single risky asset over $T = 2,520$ trading days (approximately 10 years).

2.1 Agent Types and Strategy Treatments

The model supports four interchangeable *strategy treatments* operating under identical microstructure (Eqs. 4–5). For a given experimental run, the population is partitioned as $N_{\text{BH}} = (1 - \phi)N - N_{\text{noise}}$ Buy-and-Hold (BH) agents with constant weight $w^{\text{BH}} = 1.0$, ϕN strategy-treatment agents, and $N_{\text{noise}} = 200$ noise traders with weight increments $\Delta w \sim N(0, 0.02)$ clipped to $[0, 1.5]$ providing background liquidity. The composition is well-defined as long as $\phi < 1 - N_{\text{noise}}/N = 80\%$; at the saturated adoption levels $\phi \in \{100\%\}$ all non-noise agents are strategy-treatment agents.

The strategy treatment is one of four classes:

Volatility targeting (VT). The canonical 12/VIX rule [Perchet et al., 2015],

$$w_t^{\text{VT}} = \min(12/\text{VIX}_{t-1}, 1.5), \quad (1)$$

with the lagged VIX ensuring lookahead-safe execution.

Trend-following (TF). Momentum-based weight update,

$$w_t^{\text{TF}} = \text{clip}(0.5 + s \cdot \tilde{r}_{t-1}^{(W)}, 0, 1.5), \quad (2)$$

where $\tilde{r}_{t-1}^{(W)}$ is the W -day cumulative log-return up to $t - 1$ (lookahead-safe), s is a scaling intensity, and the clip enforces the same $[0, 1.5]$ rail as VT. Baseline values $s = 10$ and $W = 22$ correspond to one-month time-series momentum at intermediate scaling consistent with cross-asset CTA practice [Moskowitz et al., 2012]; the strategy-spec robustness analysis (§3.4) sweeps $s \in \{1, 3, 5, 10\}$ and $W \in \{10, 22, 60\}$.

Mean-reversion (MR). Sign-flipped momentum,

$$w_t^{\text{MR}} = \text{clip}(0.5 - s \cdot \tilde{r}_{t-1}^{(W)}, 0, 1.5), \quad (3)$$

with the same baseline $s = 10$, $W = 22$. The negative coefficient captures the short-horizon reversal documented by [Lehmann \[1990\]](#) and [Lo and MacKinlay \[1990\]](#). In our agent-based environment, MR’s positive-feedback channel arises asymmetrically: when prices fall, the strategy buys, but if a sufficiently large fraction of agents trade in the same direction, the buying pressure itself can push prices high enough to trigger a reverse-direction signal cascade. This produces episodic price-collapse-and-recovery behavior at intermediate adoption (e.g. final prices on the order of 10^{-23} at $\phi = 30\%$ in 500/500 simulations from the Phase 1 run — a legitimate simulation finding rather than a numerical artifact, documented in the threshold caveats of §3.4).

NoiseControl (NC). Falsifiability anchor with deterministic weight $w_t^{\text{NC}} = 0.5$. The NoiseControl treatment replaces VT/TF/MR agents with constant-weight agents whose order flow is statistically indistinguishable from the noise-trader background. Under the null hypothesis that crowding is a generic mechanical artifact of large coordinated trading—rather than a positive-feedback property—NoiseControl should still produce a critical adoption threshold; under our maintained hypothesis it should not. The Phase 1 run (§3.4) confirms NoiseControl produces no detectable threshold across $\phi \in \{0\%, \dots, 100\%\}$, establishing that the cross-treatment threshold differences reflect the strategies’ positive-feedback structure rather than their headcount.

The experimental adoption variable ϕ now denotes the fraction of N agents using the *active* strategy treatment, tested at seven levels: 0%, 10%, 20%, 30%, 50%, 70%, and 100%. As ϕ rises, BH agents are replaced by treatment agents while noise traders remain fixed at $N_{\text{noise}} = 200$. This design isolates the crowding mechanism from liquidity effects: because the noise-trader population is constant across all adoption levels, any performance degradation is attributable to correlated strategy-treatment trading rather than to the mechanical reduction of background liquidity providers. Section 3.7 reports

a design validation confirming that this distinction is quantitatively important. Cross-treatment seed pairing is enforced— $\text{seed}(i) = \lfloor \phi \cdot 10^5 \rfloor + i + 42$ —so that Monte-Carlo sampling noise is paired across VT, TF, and MR comparisons within each adoption level.

2.2 Market Microstructure

Price formation follows a simplified [Kyle \[1985\]](#) model:

$$r_t = \mu + \lambda \cdot \frac{\text{NOF}_t}{N} + \varepsilon_t, \quad \varepsilon_t \sim N(0, \sigma_f), \quad (4)$$

where $\mu = 0.08/252$ is the daily drift, $\lambda = 0.005$ is the Kyle lambda (price impact per unit of normalized order flow), NOF_t is the aggregate net order flow from all agents, and $\sigma_f = 0.16/\sqrt{252}$ is the fundamental volatility. The VIX evolves endogenously:

$$\text{VIX}_t = \max\left(0, \text{VIX}_{t-1} + \kappa(\bar{V} + \gamma \cdot \Delta\sigma_t^{\text{real}} - \text{VIX}_{t-1}) + \eta_t\right), \quad (5)$$

where $\kappa = 0.03$ is the mean-reversion speed, $\bar{V} = 18$ is the long-run mean, $\Delta\sigma_t^{\text{real}} = \max(0, \hat{\sigma}_t^{20d} - 0.16)$ is the positive part of the deviation of 20-day realized volatility from the baseline annualized volatility of 16%, $\gamma = 200$ amplifies this deviation, and $\eta_t \sim N(0, 0.3)$. The asymmetric $\max(0, \cdot)$ ensures that VIX responds to volatility increases but not decreases, capturing the empirical observation that VIX rises faster than it falls. The VIX is additionally bounded to $[9, 80]$.

2.3 Feedback Structure

The model is designed to embed a positive feedback loop: when VT agents sell (because VIX rises), the negative order flow depresses prices via Eq. (4), which increases realized volatility, which raises VIX via Eq. (5), which triggers further VT selling. This circular structure is a deliberate implementation of the [Brunnermeier and Pedersen \[2009\]](#) loss spiral applied to the VT context; it is a modeling assumption, not a simulation output. The strength of this feedback loop is proportional to ϕ , and the simulation’s purpose is

to quantify the adoption threshold at which the loop becomes destabilizing.

2.4 Simulation Design

For each of the 7 adoption levels, we run $M = 500$ independent Monte Carlo simulations with different random seeds, producing a total of $3,500 \text{ simulations} \times 2,520 \text{ days} = 8.82$ million agent-day observations. Bootstrap 95% confidence intervals (2,000 replications) are reported for all key metrics. For each simulation we compute: annualized market volatility, VT strategy Sharpe ratio (annualized mean excess return over risk-free rate divided by annualized standard deviation of daily returns), maximum drawdown, excess kurtosis, skewness, flash crash frequency (daily returns exceeding 3σ), and VIX dynamics. We additionally run 200 simulations per cell across a one-at-a-time (OAT) sensitivity analysis varying λ , γ , and κ each at $\pm 50\%$ of baseline (9 parameter configurations) at three adoption levels (10%, 30%, 50%) to test threshold robustness.

3 Results

3.1 Strategy Performance Degradation

Table 1 presents the core results under fixed background liquidity ($N_{\text{noise}} = 200$). VT strategy performance exhibits a clear nonlinear pattern. At low adoption ($\phi = 10\text{--}30\%$), the VT Sharpe ratio remains near 0.47–0.50, statistically indistinguishable from the uncrowded baseline. The critical transition occurs between 50% and 70%: the Sharpe ratio falls to 0.34 at 50% (a 28% decline relative to 10%) and collapses to 0.08 at 70% (an 82% decline). At 100%, VT generates substantially negative risk-adjusted returns (-0.27).

Table 1: VT Strategy and Market Outcomes by Adoption Level

ϕ	Strategy Performance				Market Stability			
	Sharpe	95% CI	Δ Sharpe	Ret. (%)	Δ Vol (%)	Kurt.	Kurt. CI	Flash/yr
0%	—	—	—	—	baseline	−0.00	[−0.01, 0.01]	0.32
10%	0.47	[0.44, 0.50]	—	4.69	−0.1	−0.01	[−0.02, −0.00]	0.30
20%	0.50	[0.47, 0.52]	+6%	4.98	+0.8	0.00	[−0.01, 0.01]	0.31
30%	0.47	[0.44, 0.50]	−0%	4.74	+3.8	−0.00	[−0.01, 0.00]	0.31
50%	0.34	[0.31, 0.36]	−28%	3.46	+18.8	0.06	[0.05, 0.07]	0.40
70%	0.08	[0.07, 0.10]	−82%	0.90	+42.7	1.41	[1.28, 1.55]	1.09
100%	−0.27	[−0.28, −0.26]	−157%	−3.82	+119.1	61.4	[59.2, 63.4]	1.20 ^a

Results are means across 500 Monte Carlo simulations $\times$ 2,520 days each ($N = 1,000$ agents; $N_{\text{noise}} = 200$ fixed). 95% CIs from 2,000 bootstrap replications. Sharpe ratio and return are for the VT portfolio. Δ Sharpe is relative to the 10% baseline. Δ Vol is the percentage change in annualized market volatility relative to the 0% (no VT) scenario. Flash crash frequency counts daily returns $> 3\sigma$. Kurtosis is excess (Fisher); bootstrap CIs reported for all rows.

^a At 100% adoption, the moderate flash crash frequency relative to 70% reflects a measurement artifact: the standard deviation used to define the 3σ threshold is itself severely inflated (annualized vol of 35%), masking extreme events.

The nonlinearity is striking. Performance is essentially flat from 10% to 30%, then drops sharply: a linear extrapolation from the 10–30% range would predict negligible degradation at 50%, yet the actual decline is 28%. By 70%, the strategy is essentially destroyed. Figure 1 visualizes this threshold collapse with bootstrap 95% confidence bands. This pattern is consistent with a phase-transition-like threshold driven by the positive feedback structure encoded in the model (Section 2).

3.2 Market-Level Consequences

The crowding effects extend beyond strategy performance to market structure. Table 2 reports detailed market stability metrics. Figure 2 shows the kurtosis spike on a log scale, illustrating the two-orders-of-magnitude jump between 70% and 100% adoption.




Figure 1: VT tipping point: Sharpe ratio collapses non-linearly between 50–70% adoption. Markers and whiskers are means and 95% bootstrap CIs from 500 Monte Carlo simulations per adoption level. The red-shaded band marks the empirical tipping zone; the green-shaded band marks the safe zone.




Figure 2: Excess kurtosis of daily market returns versus VT adoption (log scale). Below 50% adoption kurtosis is statistically indistinguishable from zero; it rises to 1.41 at 70% and jumps roughly two orders of magnitude to 61.4 at full adoption, matching the flash-crash frequency spike in Table 1.

Table 2: Market Microstructure Effects of VT Crowding

ϕ	Ann. Vol	VIX Mean	VIX Std	Skewness	MDD	VIX Spike (%) ^b
0%	16.0%	19.4	1.76	−0.005	−33.4%	0.0
10%	16.0%	19.4	1.76	0.001	−34.2%	0.0
20%	16.1%	19.5	1.77	−0.002	−33.2%	0.0
30%	16.6%	20.0	1.89	0.002	−34.9%	0.0
50%	19.0%	22.8	2.33	−0.024	−40.2%	0.3
70%	22.8%	27.4	3.39	−0.321	−60.1%	16.2
100%	35.1%	35.0	6.45	−4.73	−91.3%	90.0

Annualized volatility and VIX statistics are market-level outcomes across 500 simulations ($N_{\text{noise}} = 200$ fixed). MDD is the mean maximum drawdown of the market portfolio.

^b Percentage of trading days where VIX exceeds the spike threshold (30).

Three patterns emerge. First, market volatility amplification is strongly nonlinear: from near-zero at 30% to +19% at 50% and +119% at 100%. Second, the VIX distribution shifts dramatically—at 70% adoption, VIX spends 16% of days in spike territory, compared to near 0% in the uncrowded baseline, and at 100% this reaches 90%. Third, return skewness turns sharply negative at high adoption levels (−0.32 at 70%, −4.7 at 100%), indicating that VT crowding manufactures left-tail events—precisely the risk that VT is designed to mitigate.

3.3 Statistical Significance

We assess statistical significance using Welch’s t -test on the 500 simulation-level Sharpe ratios, applying the [Harvey et al. \[2016\]](#) threshold of $|t| > 3.0$. The 10% versus 30% comparison yields $t = 0.05$ ($p = 0.96$), confirming that these two adoption levels are statistically indistinguishable—consistent with the “safe zone” interpretation below 30%. The degradation at 50% versus 10% yields $t = 7.12$ ($p < 0.001$), far exceeding the Harvey threshold and confirming that the tipping point between 30% and 50% represents a genuine structural break rather than simulation noise. The further collapse at 70%

versus 50% yields $t = 16.94$ ($p < 0.001$, Welch’s t -test on 500 simulation-level Sharpes per adoption level, with $\bar{S}_{50\%} = 0.336$ vs. $\bar{S}_{70\%} = 0.084$), documenting the second transition as a distinctly larger structural break than the 30–50% step.

3.4 Cross-Strategy Threshold Comparison

To establish that positive-feedback crowding is a class-level rather than VT-specific phenomenon, we compare critical adoption thresholds across the four treatments under three independently-specified detectors. The *Strict* detector requires Sharpe-drop $> 50\%$ AND excess kurtosis > 10 AND vol amplification $> 50\%$ simultaneously; the *Softer* detector loosens the kurtosis threshold to > 1 ; the *Sharpe-only* detector triggers on a Sharpe sign-flip OR a Sharpe drop exceeding 70% from the 10% baseline. The Sharpe-only specification corresponds to the criterion implicitly used by the standalone VT analysis and serves as the primary calibration anchor.

Table 3: Critical adoption thresholds across treatments under three detector specifications. Source: 10,500-simulation Phase 1 run for VT, TF, MR, NoiseControl baselines; cross-detector recompute documented in `experiments/k1262/k1262_softer_detector_table.md`.

Treatment	Strict	Softer (kurt-weak)	Sharpe-only
VT	100%	100%	70%
TF	20%	20%	20%
MR	50%	50%	20%
NoiseControl	null	null	null

“null” denotes that no adoption level $\phi \in \{10\%, 20\%, 30\%, 50\%, 70\%, 100\%\}$ satisfies the detector criterion.

All entries computed from the cell1 baseline ($\lambda = 0.005$, $\gamma = 200$, $s = 10$, $W = 22$).

Calibration. The Sharpe-only detector applied to the VT 500-MC baseline reproduces *exactly* the 70% threshold reported in the standalone analysis (Table 1)—this is the

primary anchor establishing that our cross-treatment comparisons are valid against the published headline figure.

Cross-strategy ordering. Under all three detectors, TF and MR thresholds are at or below VT. Under the Strict detector, TF crosses at 20% versus VT’s 100%; under the Sharpe-only detector, both TF and MR cross at 20% versus VT’s 70%. NoiseControl never crosses, validating the falsifiability anchor: a deterministic-weight strategy class with no positive-feedback channel does *not* produce a threshold even at 100% adoption.

Strategy-spec robustness. A 16,800-simulation Phase 2 robustness sweep extends this comparison to 12 (scaling, window) cells covering $s \in \{1, 3, 5, 10\}$ and $W \in \{10, 22, 60\}$. Under the softer detector, TF threshold $<$ VT threshold in 12/12 cells; MR threshold $\leq$ VT in 12/12 cells (Table 4).

Table 4: TF / MR critical adoption (softer detector) across scaling $\times$ window cells; VT cell1 reference is 100% under the softer detector. Source: `experiments/k1262/k1262_threshold_matrix.md` (Phase 2, 16,800 simulations).

Scaling $s \setminus$ Window W	10 days	22 days	60 days
1	TF: 70% / MR: 70%	TF: 70% / MR: 70%	TF: 70% / MR: 70%
3	TF: 30% / MR: 50%	TF: 30% / MR: 50%	TF: 30% / MR: 50%
5	TF: 30% / MR: 50%	TF: 20% / MR: 50%	TF: 20% / MR: 50%
10	TF: 20% / MR: 50%	TF: 20% / MR: 50%	TF: 20% / MR: 50%

The TF threshold ranges from 20% (aggressive scaling) to 70% (mild scaling) but never exceeds VT’s 100% reference under the softer detector; MR ranges between 50% and 70%. The qualitative ordering $\text{TF/MR} \leq \text{VT}$ is therefore robust to both signal scaling intensity and signal window choice—i.e. the family-level finding does not depend on a single strategy parameterization. Section 3.6 extends this robustness analysis to the underlying microstructure parameters λ and γ .

3.5 The Feedback Structure

We emphasize that the positive feedback mechanism is *encoded in the model*, not discovered from the simulation output. The VIX responds to realized volatility (Eq. 5), VT agents respond to VIX, and their trading affects prices (Eq. 4), which in turn affects realized volatility. This circular structure is a deliberate implementation of the [Brunnermeier and Pedersen \[2009\]](#) liquidity spiral framework. What the simulation *does* reveal is the *quantitative threshold behavior*: the feedback loop’s consequences are negligible below 30% adoption, become material at 50%, and are catastrophic at 70%. Making the feedback explicit, the aggregate VT order flow on day t is $\phi N \Delta w_t^{\text{VT}}$, where the VT weight change uses the capped rule from Eq. (1),

$$\Delta w_t^{\text{VT}} = \min\left(\frac{12}{\text{VIX}_{t-1}}, 1.5\right) - \min\left(\frac{12}{\text{VIX}_{t-2}}, 1.5\right). \quad (6)$$

Its contribution to the price via Eq. (4) is $\lambda \phi \Delta w_t^{\text{VT}}$, which is *linear* in adoption ϕ ; the *nonlinearity* observed in our results arises because this order flow itself amplifies subsequent VIX increases via Eq. (5), closing the loop endogenously. The simulation thus serves as a computational laboratory for quantifying a mechanism whose qualitative direction is known, but whose critical mass has not previously been estimated.

3.6 Sensitivity Analysis

To assess whether the 50–70% tipping point is an artifact of our baseline calibration, we vary λ (market impact), γ (VIX sensitivity to realized vol), and κ (VIX mean-reversion speed) each by $\pm 50\%$, producing 9 one-at-a-time (OAT) parameter variations tested at $\phi \in \{10\%, 30\%, 50\%\}$ with 200 simulations per cell (Table 5).

Table 5: Sensitivity of VT Tipping Point to Key Parameters (Fixed Liquidity)

Parameter	Value	Sharpe ₁₀	Sharpe ₃₀	Sharpe ₅₀	Threshold ^c
λ	0.0025 (−50%)	0.52	0.49	0.43	>50%
	0.005 (baseline)	0.49	0.45	0.35	>50%
	0.0075 (+50%)	0.52	0.40	0.23	30–50%
γ	100 (−50%)	0.47	0.51	0.36	>50%
	200 (baseline)	0.52	0.47	0.32	>50%
	300 (+50%)	0.48	0.41	0.31	>50%
κ	0.015 (−50%)	0.45	0.48	0.32	>50%
	0.03 (baseline)	0.49	0.47	0.34	>50%
	0.045 (+50%)	0.51	0.44	0.32	>50%

Each row varies one parameter while holding others at baseline. 200 simulations per cell ($N_{\text{noise}} = 200$ fixed).

^c Threshold is the adoption region where Sharpe degradation first exceeds 50% relative to the 10% baseline. Under fixed liquidity, only high- λ produces a threshold below 50%. λ : Kyle market impact; γ : VIX sensitivity to realized volatility; κ : VIX mean-reversion speed.

The tipping point is *conditionally* stable. Under fixed background liquidity, eight of nine parameter combinations produce a threshold above 50%; only the high- λ scenario (+50%) pulls the threshold into the 30–50% region. VIX dynamics (γ , κ) have limited influence: all six γ/κ variants produce a threshold above 50%. Market impact (λ) remains the dominant driver: the Sharpe spread across λ variants at $\phi = 50\%$ is 0.21 (0.43 to 0.23), compared to 0.05 for γ and 0.02 for κ . This is economically intuitive—the feedback loop’s severity depends on how much correlated selling moves prices, not on how VIX tracks volatility.

3.7 Design Validation: Fixed vs. Scaled Liquidity

An earlier version of this model scaled noise traders proportionally with $(1 - \phi)$, so that at high VT adoption the number of background liquidity providers fell dramatically (from 300 to near zero at $\phi = 100\%$). This confounded two effects: crowding (correlated VT trading) and liquidity evaporation (fewer noise traders to absorb order flow). We address this by fixing $N_{\text{noise}} = 200$ across all adoption levels.

The comparison is revealing. When noise traders are scaled with VT adoption, the threshold appears at 30–50% (Sharpe collapses from 0.43 to 0.18 between 30% and 50%). Fixing background liquidity shifts the threshold to 50–70% (Sharpe falls from 0.47 to 0.34 at 50%, and only collapses to 0.08 at 70%). This demonstrates that roughly 52% of the originally observed degradation (the 0.13 Sharpe-drop under fixed liquidity relative to the 0.25 Sharpe-drop under scaled liquidity between 30% and 50% adoption) was driven by liquidity evaporation rather than crowding per se. The fixed-liquidity design is more conservative—it attributes less degradation to crowding—and we regard these as the more reliable estimates.

4 Discussion

4.1 Policy Implications

Regulators should monitor VT adoption rates as a systemic risk indicator. The 50–70% tipping point—while far from current levels below 5%—provides a quantitative benchmark. The ESRB and FSOC could incorporate VT adoption metrics into macroprudential surveillance, alongside existing crowding measures [Baltas, 2019].

4.2 Implications for Practitioners

At current adoption levels, VT remains viable. However, the nonlinear tipping point means degradation could accelerate rapidly once critical mass is reached. Practitioners should monitor proxy indicators: the correlation between VIX changes and equity fund

flows, volatility-linked AUM concentration, and anomalous intraday reversals during VIX spikes.

4.3 Limitations

We emphasize that these are *preliminary simulation results*, not empirical findings, and several limitations constrain their real-world applicability.

First, our model uses a constant λ , whereas Kyle [1985] derives λ endogenously from the information structure. In real markets, liquidity evaporates during stress as market makers widen spreads, which would amplify the feedback loop and potentially lower the tipping point. Conversely, in normal times, deep liquidity would moderate price impact. Our sensitivity analysis (Table 5) confirms that λ is the most influential parameter, suggesting that endogenizing it is the highest-priority extension.

Second, agents within each class are homogeneous: all VT agents use identical 12/VIX rules with the same rebalancing frequency. In practice, VT strategies vary widely in their volatility estimators, target levels, rebalancing frequencies, and position limits. This heterogeneity would likely smooth the tipping point and shift it to higher adoption levels.

Third, the model excludes transaction costs, margin constraints, and short-selling restrictions, all of which would moderate extreme outcomes. It also lacks adaptive learning—agents do not switch strategies in response to poor performance, whereas real investors would eventually exit a degraded strategy.

Fourth, our simplified VIX is a realized-volatility tracker that omits the variance risk premium embedded in the true CBOE VIX (option-implied volatility). This distinction may affect the magnitude of the feedback—the real VIX typically exceeds realized volatility, so VT agents in practice hold lower weights than our model implies—but does not alter the direction of the crowding mechanism.

Fifth, with $N = 1,000$ agents, the model is small relative to real markets with millions of participants. The price impact per agent is correspondingly larger than in reality.

Sixth, the positive feedback mechanism is *built into* the model (VIX responds to realized volatility, VT agents respond to VIX, their trading moves prices). We do not

claim to discover this feedback—it is a structural feature of the simulation—but rather to quantify the adoption level at which it becomes destabilizing.

These limitations suggest that our quantitative thresholds (50–70%) should be interpreted as order-of-magnitude estimates rather than precise critical values. The design validation (Section 3.7) demonstrates that even the *direction* of the threshold depends on modeling choices—an earlier design with scaled liquidity produced a threshold 20 percentage points lower. The emerging empirical literature on VT crowding, including ongoing work on realized VT fund flows, will be essential for calibrating these thresholds to real markets. The qualitative finding—that VT strategies exhibit a nonlinear crowding tipping point whose severity depends on market impact—is more robust than any specific numerical threshold.

5 Conclusion

This paper presents preliminary agent-based evidence that volatility targeting strategies exhibit a nonlinear crowding tipping point. Using a Kyle (1985) market microstructure with 1,000 heterogeneous agents, 200 fixed noise traders, and 500 Monte Carlo simulations, we find that VT performance remains stable below 30% adoption but degrades sharply between 50% and 70%. The model encodes a Brunnermeier–Pedersen positive feedback structure, and we quantify—rather than discover—its threshold behavior under controlled conditions. At 70% adoption, the strategy is essentially destroyed (Sharpe 0.08) and market kurtosis rises to 1.4; at 100%, kurtosis reaches 61. A design validation confirms that approximately half the degradation in a naïve specification is attributable to liquidity evaporation rather than crowding. Sensitivity analysis across 9 one-at-a-time (OAT) parameter variations confirms that the threshold is conditionally stable—shifting primarily with market impact (λ) while remaining robust to VIX dynamics.

Current VT adoption appears safely below this threshold. However, the nonlinear nature of the tipping point—where small increases in adoption produce disproportion-

ately large degradation—suggests that monitoring VT growth is warranted. Future work should extend the model with endogenous λ , heterogeneous VT specifications, adaptive strategy switching, and calibration to real market microstructure data to refine the threshold estimate.

References

- Baltas, N. (2019). The impact of crowding in alternative risk premia investing. *Financial Analysts Journal*, 75(3), 89–104.
- Barroso, P. and Detzel, A. (2021). Do limits to arbitrage explain the benefits of volatility-managed portfolios? *Journal of Financial Economics*, 140(3), 744–767. <https://doi.org/10.1016/j.jfineco.2021.02.009>
- Bookstaber, R., Paddrik, M., and Tivnan, B. (2014). An agent-based model for financial vulnerability. OFR Working Paper 14-05, Office of Financial Research.
- Brunnermeier, M. K. and Pedersen, L. H. (2009). Market liquidity and funding liquidity. *Review of Financial Studies*, 22(6), 2201–2238. <https://doi.org/10.1093/rfs/hhn098>
- Cederburg, S., O’Doherty, M. S., Wang, F., and Yan, X. S. (2020). On the performance of volatility-managed portfolios. *Journal of Financial Economics*, 138(1), 95–117. <https://doi.org/10.1016/j.jfineco.2020.04.015>
- European Central Bank. (2020). Volatility-targeting strategies and the market sell-off. *Financial Stability Review*, May.
- Gennotte, G. and Leland, H. E. (1990). Market liquidity, hedging, and crashes. *American Economic Review*, 80(5), 999–1021.
- Harvey, C. R., Liu, Y., and Zhu, H. (2016). ...and the cross-section of expected returns. *Review of Financial Studies*, 29(1), 5–68. <https://doi.org/10.1093/rfs/hhv059>

- Harvey, C. R., Hoyle, E., Korgaonkar, R., Rattray, S., Sargaison, M., and Van Hemert, O. (2018). The impact of volatility targeting. *Journal of Portfolio Management*, 45(1), 14–33.
- Kyle, A. S. (1985). Continuous auctions and insider trading. *Econometrica*, 53(6), 1315–1335.
- Moreira, A. and Muir, T. (2017). Volatility-managed portfolios. *Journal of Finance*, 72(4), 1611–1644. <https://doi.org/10.1111/jofi.12513>
- Perchet, R., de Carvalho, R. L., Heckel, T., and Moulin, P. (2015). Predicting the success of volatility targeting strategies: Application to equities and other asset classes. *Journal of Alternative Investments*, 18(3), 21–38.
- Cole, C. (2017). Volatility and the alchemy of risk. Artemis Capital Management, Research Report.
- Danielsson, J., Shin, H. S., and Zigrand, J.-P. (2012). Procyclical leverage and endogenous risk. Working Paper, London School of Economics.
- LeBaron, B. (2006). Agent-based computational finance. In L. Tesfatsion and K. L. Judd (Eds.), *Handbook of Computational Economics*, Vol. 2, pp. 1187–1233. Elsevier.
- Liu, F., Tang, X., and Zhou, G. (2019). Volatility-managed portfolio: Does it really work? *Journal of Portfolio Management*, 46(1), 38–51. <https://doi.org/10.3905/jpm.2019.1.107>
- Moskowitz, T. J., Ooi, Y. H., and Pedersen, L. H. (2012). Time series momentum. *Journal of Financial Economics*, 104(2), 228–250. <https://doi.org/10.1016/j.jfineco.2011.11.003>
- Asness, C. S., Moskowitz, T. J., and Pedersen, L. H. (2013). Value and momentum everywhere. *Journal of Finance*, 68(3), 929–985. <https://doi.org/10.1111/jofi.12021>

- Lehmann, B. N. (1990). Fads, martingales, and market efficiency. *Quarterly Journal of Economics*, 105(1), 1–28. <https://doi.org/10.2307/2937816>
- Lo, A. W. and MacKinlay, A. C. (1990). When are contrarian profits due to stock market overreaction? *Review of Financial Studies*, 3(2), 175–205. <https://doi.org/10.1093/rfs/3.2.175>
- Cont, R. and Bouchaud, J.-P. (2000). Herd behavior and aggregate fluctuations in financial markets. *Macroeconomic Dynamics*, 4(2), 170–196. <https://doi.org/10.1017/S1365100500015029>